

# Day 04

---


## The Modeling Project

MSU REU Machine Learning Short Course

# The Full ML Loop

---

You've seen each piece. Today you run the whole thing yourself.




**Understand  
the Problem**

## Today's Dataset — Molecular Atomization Energy

---

Predict the energy required to pull a molecule apart, from its atomic structure.



**1275 features** from the Coulomb matrix representation of each molecule. Most of them are noise.

# Start with a terrible model

---

Build the simplest possible model first.

```
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)


model = LinearRegression()
model.fit(X_train, y_train)
print(r2_score(y_test, model.predict(X_test)))    # probably bad
```

Then explore the data to understand **why** it's bad.

**Hint:** plot a histogram of each feature. What do you notice?

## Feature selection — less is more

Not all 1275 features carry signal. Recursive Feature Elimination finds which ones to keep.



# Cross-validation — how confident are you?

---

One train/test split gives you one number. Run the model many times to get a **distribution**.


Each fold takes a turn as the test set.  
This gives 5 independent estimates of model performance.

|       |       |       |       |       |
|-------|-------|-------|-------|-------|
| TEST  | train | train | train | train |
| train | TEST  | train | train | train |
| train | train | TEST  | train | train |
| train | train | train | TEST  | train |
| train | train | train | train | TEST  |

# Dimensionality reduction — PCA

---

1275 correlated features → a small number of uncorrelated components that capture most of the variance.



# Today's Activity

---

Work through **Activity 04: Modeling Project** — open-ended by design.

1. Load the molecular energy dataset
2. Build an initial model — expect it to be bad
3. Explore: why is it bad? Remove useless features.
4. Re-evaluate. Does it improve?
5. Implement cross-validation to estimate performance with uncertainty.

**Notes:** [Cross-Validation](#) · [RFE](#) · [PCA](#)

Use everything from Days 1–3. Ask questions. Improve iteratively.