

Day 02: Classification with KNN

using `scikit-learn`

MSU REU Machine Learning Short Course

Learning Goals for Today

By the end of this activity, you'll be able to:

1. **Describe the ML pipeline**: this is the same structure you'll reuse for most models
2. **Split and scale data** correctly, and explain **why** each step matters
3. **Build and train a KNN classifier** using `scikit-learn`
4. **Evaluate a model** with precision, recall, F1, and the confusion matrix
5. **Improve a model** by adding features and tuning hyperparameters

This is the workflow professional data scientists use every day.

From Exploration to Modeling


Last week: **which features might separate classes?**

Today: we build a model that makes that separation automatic.

	Day 01	Day 02
Goal	Understand the data	Predict the class
Tools	<code>pandas</code> , <code>seaborn</code>	<code>scikit-learn</code>
Output	Plots and intuitions	A trained classifier


Same dataset. New question: can a machine learn the rules?

The ML Pipeline



Every model you build in class follows this same structure — for classification **and** regression.

Train / Test Split



Train / Test Split

Why? You can't evaluate yourself on data you learned from.

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

- 80% trains the model
- 20% evaluates it — **unseen data**
- `random_state=42` makes splits reproducible

Feature Scaling

Why? KNN uses distance.

Redshift ranges 0–7, photometry 10–33. Without scaling, redshift dominates every distance calculation (smaller distances are closer neighbors).

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```


Fit on training data only — then apply the same transformation to test.

Why Rescale?

Feature	Unscaled σ	Scaled σ
Photometry u, g, z	$\sim 31.7 \leftarrow$ large	1.0
Photometry i, z	1.8 $\leftarrow$ small	1.0
Redshift	0.73 $\leftarrow$ small	1.0

KNN — You Are Your Neighbors

To classify a new point, find the **k** closest training points and take a majority vote.



Building the Model

The `scikit-learn` pattern is nearly the same for every algorithm.

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix

knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train_scaled, y_train)      # learn from training data

y_pred = knn.predict(X_test_scaled)   # predict on unseen data
print(classification_report(y_test, y_pred))
```

Fit → Predict → Evaluate. This four-step pattern carries forward to every model we build.

Evaluation — The Confusion Matrix

True Positive (TP)
Said YES · Was YES ✓

False Positive (FP)
Said YES · Was NO ✗

False Negative (FN)
Said NO · Was YES ✗

True Negative (TN)
Said NO · Was NO ✓

Evaluation — Precision, Recall, F1

- **Precision** — when I predict QSO, how often am I right?

$$\text{Precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

- **Recall** — of all real QSOs, how many did I catch?

$$\text{Recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

- **F1** — harmonic mean; use it when classes are imbalanced (which is common)


$$\text{F1} = 2 \cdot \frac{\text{Precision} \cdot \text{Recall}}{\text{Precision} + \text{Recall}}$$

Good values for all three? **Between 0.8 and 1.0**

But the "best" value depends on your goals — there's no single right answer.


Redshift Changes Everything

Without redshift



86% accuracy

With redshift




>99% accuracy

Why does one measurement change everything?

Tuning k — Which Value is Best?

Performance varies with the number of neighbors. Finding the best k is part of the job.



Today's Activity

Work through **Activity 02: Classification with** `scikit-learn`

1. Build a **2-class KNN** (STAR vs QSO) using spectral features only
2. **Tune k** — find the value that maximizes F1 score
3. Add **redshift** — compare the confusion matrices before and after
4. Extend to the **3-class problem** (STAR, GALAXY, QSO)

Notes: [Methods & Validation](#) · [Support Vector Machines](#)

Because some of these ideas might be new, there's a lot of scaffolding in the notebook. Don't worry about memorizing it.

Understand the workflow and how to use the tools. You'll reuse this pattern for every model you build.